

Inference Latency Comparison: Sparse vs Dense Attention (Lower is Better)

